

ASSIGNMENT 4
SHIVAM AGNIHOTRI (SA22BB)

1. All your code (all working) that you wrote to train, validate and test the model.

(complete code in folder)

training

```
while epoch < max_epochs:
    epoch += 1
    print(f'epoch no: {epoch}')
    train_loss = 0.0
    valid_loss = 0.0
    accuracy_val = 0.0
```

```
model.train()
```

training

```
for index, batch in enumerate(tqdm(train_dl)):
    optimizer.zero_grad()
    # get the input from batch
    input_data, target = batch
    # pass the input to model and get prediction
    output = model(input_data)
    # calculate the loss and store it in variable loss using the criterion
    loss = F.cross_entropy(output, target)
    # backpropagate the loss
    loss.backward()
    train_loss = train_loss + ((1 / (index + 1)) * (loss - train_loss))
```

```
optimizer.step()
scheduler.step()
```

```
model.eval()
```

validation

```
with torch.no_grad():
```

```

for index, batch in enumerate(tqdm(test_dl)):

    # get the input from batch
    input_data, target = batch
    # pass the input to model and get prediction
    output = model(input_data)
    # calculate the loss and store it in variable loss using the criterion
    loss = F.cross_entropy(output, target)
    # calculate the accuracy and store it in variable acc using the function called accuracy
    acc = accuracy(output, target)
    valid_loss = valid_loss + ((1 / (index + 1)) * (loss - valid_loss))
    accuracy_val = accuracy_val + ((1 / (index + 1)) * (acc - accuracy_val))

print(f"\ntrain_loss: {train_loss:.2f} \n"
      f"valid_loss: {valid_loss:.2f} \n"
      f"acc: {accuracy_val:.2f}")

writer.add_scalar("train_loss", train_loss, global_step=epoch)
writer.add_scalar("valid_loss", valid_loss, global_step=epoch)
writer.add_scalar("accuracy", accuracy_val, global_step=epoch)

if epoch % config["save_epoch"] == 0:
    model_save(Path(config["model_save_path"]), model, optimizer, scheduler, epoch,
"current")

    if accuracy_val > max_acc:
        max_acc = accuracy_val
        model_save(Path(config["model_save_path"]), model, optimizer, scheduler, epoch,
"best")

writer.flush()

writer.close()

```

2. Screenshot and steps of training and validating the model.

```
train_loss: 0.14
valid_loss: 0.10
acc: 0.97
```

The model is trained and validated using the following steps:

- The code initializes the necessary variables, such as the epoch number, maximum accuracy, and maximum number of epochs specified in the configuration file.
- The model is put into training mode using `model.train()`.

```
97     # training
98     while epoch < max_epochs:
99         epoch += 1
100        print(f"epoch no: {epoch}")
101        train_loss = 0.0
102        valid_loss = 0.0
103        accuracy_val = 0.0
104
105        model.train()
```

- The training dataset is loaded into a `DataLoader`, which returns a batch of images and their corresponding labels for each iteration of the loop.
- For each batch, the optimizer's gradients are zeroed using `optimizer.zero_grad()`.
- The model is fed the input data (images) using `output = model(input_data)`. The output is the predicted class labels for each image in the batch.

```
106
107     # training
108     for index, batch in enumerate(tqdm(train_dl)):
109         optimizer.zero_grad()
110         # get the input from batch
111         input_data, target = batch
112         # pass the input to model and get prediction
113         output = model(input_data)
```

- The model's output is compared to the actual labels for each image in the batch, and the loss is calculated using cross-entropy loss (`F.cross_entropy(output, target)`).
- The loss is then backpropagated through the network using `loss.backward()`.
- The training loss is updated using the formula `train_loss = train_loss + ((1 / (index + 1)) * (loss - train_loss))`. This formula is used to calculate a running average of the training loss.

```
114         # calculate the loss and store it in variable loss using the criterion
115         loss = F.cross_entropy(output, target)
116         # backpropagate the loss
117         loss.backward()
118         train_loss = train_loss + ((1 / (index + 1)) * (loss - train_loss))
```

- The optimizer's weights are updated using `optimizer.step()`.

- After all batches in the training dataset have been iterated over, the learning rate scheduler's step function is called using `scheduler.step()`.
- The model is put into evaluation mode using `model.eval()`.

```
120         optimizer.step()
121         scheduler.step()
122
123         model.eval()
```

- The testing dataset is loaded into a `DataLoader`, which returns a batch of images and their corresponding labels for each iteration of the loop.
- For each batch, the model's output is computed using `output = model(input_data)`.
- The loss is calculated using cross-entropy loss (`F.cross_entropy(output, target)`), and the accuracy is calculated using a custom function `accuracy(output, target)`.

```
124         # validation
125         with torch.no_grad():
126             for index, batch in enumerate(tqdm(test_dl)):
127                 # get the input from batch
128                 input_data, target = batch
129                 # pass the input to model and get prediction
130                 output = model(input_data)
131                 # calculate the loss and store it in variable loss using the criterion
132                 loss = F.cross_entropy(output, target)
```

- The validation loss and accuracy are updated using the same running average formula used for the training loss.
- After all batches in the testing dataset have been iterated over, the training and validation losses and accuracy are printed.
- The current epoch's losses and accuracy are written to TensorBoard using the `SummaryWriter` object.

```
133         # calculate the accuracy and store it in variable acc using the function called accuracy
134         acc = accuracy(output, target)
135         valid_loss = valid_loss + ((1 / (index + 1)) * (loss - valid_loss))
136         accuracy_val = accuracy_val + ((1 / (index + 1)) * (acc - accuracy_val))
137
138         print(f"\ntrain_loss: {train_loss:.2f} \n"
139               f"valid_loss: {valid_loss:.2f} \n"
140               f"acc: {accuracy_val:.2f}")
141
142         writer.add_scalar("train_loss", train_loss, global_step=epoch)
143         writer.add_scalar("valid_loss", valid_loss, global_step=epoch)
144         writer.add_scalar("accuracy", accuracy_val, global_step=epoch)
```

- If the current epoch is a multiple of the `save_epoch` parameter specified in the configuration file, the model is saved to disk using the `model_save()` function. If the current accuracy is greater than the current maximum accuracy, the model is saved as the best checkpoint using the same function.
- The loop continues until the maximum number of epochs is reached, at which point the `SummaryWriter` object is flushed and closed.

```

145
146         if epoch % config["save_epoch"] == 0:
147             model_save(Path(config["model_save_path"]), model, optimizer, scheduler, epoch, "current")
148             if accuracy_val > max_acc:
149                 max_acc = accuracy_val
150                 model_save(Path(config["model_save_path"]), model, optimizer, scheduler, epoch, "best")
151
152     writer.flush()
153     writer.close()
154
155

```

Run screenshots:

[illegible]

